




Over-Quantization: Investigating Limits and Error Propagation

Guided by:
Dr. M. J. Camel Mary Belinda

Presented by:
D. Mani Kumar Reddy (192325095)
M. Siddartha(192311314)
M. Pranay(192525382)
R. Dinesh (192421117)
P. Likith Reddy (192425134)

The Quantization Foundation



Process

Mapping continuous signals to finite discrete sets (levels or integers). Core trade-off: precision vs. storage/computation.




6 dB Rule

Each additional bit $\approx +6.02$ dB in SQNR; practical gains diminish with perceptual masking and system noise floors.



Key Assumption


Quantization noise modelled as uniform, white, and uncorrelated with input — a simplification that breaks in low-bit or structured systems.



When Assumptions Fail

Many theoretical results assume many quantization bins and small step sizes. In practical low-bit regimes, feedback, predictive coding, and modulators introduce correlation between signal and quantization error, producing patterned distortion rather than white noise.

Visualising Information Loss



Low bit-depth creates visible banding (false contouring) as continuous gradations collapse into discrete steps. Reductions in intensity levels force sharp boundaries, losing subtle detail and texture.

Gamut mapping for limited displays/printers further shifts hues and compresses shadow detail, causing detail loss near black points and colour clipping.

The Audio & Signal Threshold

Perceptual Sensitivity

Human hearing is logarithmic; companding (μ -law, A-law) allocates bits to preserve low-level detail.


Dynamic Range

16-bit audio (~ 98 dB SQNR) usually exceeds environmental noise floors, but processing and resampling can expose quantization artefacts.

Clipping vs Distortion

Analog systems compress gradually; digital clipping produces hard bounds and harmonically-rich distortion that is perceptually intrusive.






Limitation: Layer-Wise Error Accumulation

Standard post-training quantization treats layers independently. Errors introduced early propagate downstream, interacting with non-linear activations and attention mechanisms — resulting in amplified, structured degradation. In deep transformers, this accumulation often grows near-exponentially.

The Scale of Modern Failure

At INT4/INT3, LLMs show severe accuracy loss; INT2 commonly diverges. Layer-wise calibration often overfits to calibration data, reducing robustness on real inputs. These regimes reveal that naive quantization is not robust for modern large models.



The Solution: Explicit Error Propagation (QEP)



Framework

Quantization Error Propagation (QEP) explicitly tracks upstream errors and inserts corrective terms during layer calibration.



Tunable Strength

Adjustable compensation vs compute trade-off — tuneable to preserve throughput while reducing accumulated distortion.



Compatibility

Designed to integrate with GPTQ, AWQ, QuIP and other PTQ pipelines without disrupting end-to-end deployment.

Prototyping and Performance

Visual demos comparing baseline quantised models versus QEP-enhanced models show reduced error growth and clearer outputs. Benchmarks indicate QEP restores usable accuracy in INT2 regimes where baselines fail.

- Restores task accuracy at very low bit-widths
- Correction terms are computationally lightweight
- Robust across calibration sets — less overfitting





Conclusion: Moving Beyond Uniformity

Recognise Loss

Quantization is inherently lossy — assume structure in errors rather than uniform white noise.

Holistic Calibration

Shift from independent layer optimisation to propagation-aware, system-level calibration.

Future Directions

Adaptive bit-depths, dynamic propagation control, and hybrid correction schemes are promising paths to efficient, high-fidelity deployment.